

Comparing priors over binary matrices within latent feature models framework

Supervisor: Lorenzo Ghilotti

V. Iapaolo, L. Maci, Z. Mehboob, V. Montefusco,
G.T. Morra, A. Napolitano

Department of Mathematical Engineering
Politecnico di Milano

Group project, 9 January 2024

The Beta-Bernoulli model: Gibbs Sampler

The aim is to sample binary matrices Z according to the Beta-Bernoulli Model.

Let $A \in \mathbb{R}^{\tilde{N} \times D}$, $X \in \mathbb{R}^{N \times D}$, $\alpha < 0$, $\theta > -\alpha$:

Algorithm 1: GibbsSamplerbetabernoulli

Data: $X, A, \sigma_X, \sigma_A, \text{BurnInIterations}, \text{Iterations}$

Result: collection of Z matrices

for $iter \leftarrow 1$ **to** $\text{BurnInIterations} + \text{Iterations}$ **do**

for $i \leftarrow 1$ **to** N **do**

$K \leftarrow$ number of non-empty columns of Z_{-i} ;

for $j \leftarrow 1$ **to** K **do**

 Sample already observed features;

end

 Sample new features;

end

 Compute $E[A|X, Z]$ and $\log[P(X, Z)] \leftarrow \text{Eq. (12)} + (21)$

if $iter > \text{BurnInIterations}$ **then**

 Add Z to the collection;

end

end

Sampling observed features:update $z_{i,j}$

Let K be the number of non-null columns of Z_{-i} . Let $m \in \mathbb{R}^K$ be the vector that tells us how many subjects (except i) have feature j , $j \leq K$

Algorithm 2: Sample observed features

Data: $X, A, Z, \sigma_x, \sigma_A, i, j, m_j$

Result: Z

$z_{i,j} \leftarrow 1;$

$P(X|Z) \leftarrow$ see eq. (21);

$P(z_{ij} = 1 | z_{-i}) \leftarrow \frac{m_j - \alpha}{\theta + (n-1)};$

$z_{i,j} \leftarrow 0;$

$P(X|Z) \leftarrow$ see eq. (21);

$P(z_{ij} = 0 | z_{-i}) \leftarrow 1 - P(z_{ij} = 1 | z_{-i});$

$P(z_{il} = 1 | X, Z_{-i,j}) \propto P(X|Z) \cdot P(z_{ij} = 1 | z_{-i});$

$P(z_{il} = 0 | X, Z_{-i,j}) \propto P(X|Z) \cdot P(z_{ij} = 0 | z_{-i});$

Normalize the probabilities;

$prob_param \leftarrow P(z_{il} = 1 | X, Z_{-i,j});$

$z_{ij} \sim \text{Bernoulli}(prob_param)$

Sampling of new features

Let w be the number of new features observed by subject i . In this step we sample w :

Algorithm 3: Sample new features

Data: $X, A, Z, \sigma_x, \sigma_A, i$

Result: Z

$W \leftarrow \tilde{N} - K;$

$prob_param = 1 - \frac{\theta + \alpha + n - 1}{\theta + n - 1};$

Initialize **prob_vector** of dimension W to 0;

for $h \leftarrow 0, 1, \dots, W$ **do**

$prob_vector_h \leftarrow \text{BinomialProbability}(W, prob_param, K = h);$

 update Z ;

 Calculate $P(X|Z)$;

$P(w = h | X, Z) \propto P(X|Z) \cdot prob_vector_h;$

end

Normalize **prob_vector**;

Sample $w \sim \text{Discrete}(\text{prob_vector});$

Update Z ;

The Indian Buffet Process: Gibbs Sampler

The aim is to sample binary matrices Z according to the IBP.

Let $A \in \mathbb{R}^{K \times D}$, $X \in \mathbb{R}^{N \times D}$ (K unbounded), $0 \leq \alpha \leq 1, \theta \geq \alpha, \gamma \geq 0$:

Algorithm 4: GibbsSamplerIBP

Data: $X, A, \sigma_X, \sigma_A, \text{BurnInIterations}, \text{Iterations}$

Result: collection of Z matrices

```
for  $iter \leftarrow 1$  to  $\text{BurnInIterations} + \text{Iterations}$  do
  for  $i \leftarrow 1$  to  $N$  do
     $K \leftarrow$  number of non-empty columns of  $Z_{-i}$ ;
    for  $j \leftarrow 1$  to  $K$  do
      | Sample already observed features;
    end
    | Sample new features, IBP;
  end
  Compute  $E[A|X, Z]$  and  $\log[P(X, Z)] \leftarrow \text{Eq. (12)} + (21)$ 
  if  $iter > \text{BurnInIterations}$  then
    | Add  $Z$  to the collection;
  end
end
```

IBP: Sampling new features

Let w be the number of new features observed by subject i . In this step we sample w :

Algorithm 5: Sample new features, IBP

Data: $X, A, Z, \sigma_x, \sigma_A, \alpha, \theta, \gamma$

Result: Z

$W \leftarrow UB - K$; // $UB = 10$ if $K > 20$; 5 otherwise

$prob_param = \gamma \cdot \frac{(\theta + \alpha)_n}{(\theta + 1)_n}$;

Initialize **prob_vector** of dimension W to all 0;

for $h \leftarrow 0, 1, \dots, W$ **do**

$prob_vector_h \leftarrow \text{PoissonProbability}(prob_param, k = h)$;

 Update Z ;

 Calculate $P(X|Z)$;

$P(w = h | X, Z) \propto P(X|Z) \cdot prob_vector_h$;

end

Normalize **prob_vector**;

Sample $w \sim \text{Discrete}(\text{prob_vector})$;

Update Z ;

Metropolis-Hastings Step for σ_x

Algorithm 6: Metropolis-Hastings step for updating σ_x

Data: $Z, X, A, \sigma_A, \sigma_x^{(\text{current})}$, proposal variance, prior variance

Result: Updated σ_x

Generate proposal $\sigma_x^{(\text{new})}$ from $\mathcal{N}(\sigma_x^{(\text{current})}, \sqrt{\text{proposal variance}})$;

if $\sigma_x^{(\text{new})} \leq 0$ **then**

 | return $\sigma_x^{(\text{current})}$; ; // Ensure proposed σ_x is positive

end

Compute $M^{(\text{current})}$ and $M^{(\text{new})}$ matrices;

Calculate log-likelihoods, $\log P(X|Z, \sigma_x^{(\text{current})}, \sigma_A)$, with $\sigma_x^{(\text{current})}$;

Calculate log-likelihoods, $\log P(X|Z, \sigma_x^{(\text{new})}, \sigma_A)$, with $\sigma_x^{(\text{new})}$;

Calculate log-priors, $\log \left(e^{-\frac{1}{2} \frac{\sigma_x^{(\text{current})^2}}{\text{prior variance}}} \right)$, for current $\sigma_x^{(\text{current})}$;

Calculate log-priors $\log \left(e^{-\frac{1}{2} \frac{\sigma_x^{(\text{new})^2}}{\text{prior variance}}} \right)$ for new $\sigma_x^{(\text{new})}$;

Compute log acceptance ratio $\log \left(\frac{\text{new_log_likelihood} \times \text{new_log_prior}}{\text{current_log_likelihood} \times \text{current_log_prior}} \right)$;

Sample from uniform distribution;

if $\log(\text{uniform}()) < \log(\text{acceptance ratio})$ **then**

 | return $\sigma_x^{(\text{new})}$;

else

 | return $\sigma_x^{(\text{current})}$;

end

Metropolis-Hastings Step for σ_a

Algorithm 7: Metropolis-Hastings step for updating σ_a

Data: $Z, X, A, \sigma_x, \sigma_a^{(\text{current})}$, proposal variance, prior variance

Result: Updated σ_a

Generate proposal $\sigma_a^{(\text{new})}$ from $\mathcal{N}(\sigma_a^{(\text{current})}, \sqrt{\text{proposal variance}})$;

if $\sigma_a^{(\text{new})} \leq 0$ **then**

 | return $\sigma_a^{(\text{current})}$; ; // Ensure proposed σ_a is positive

end

Compute $M^{(\text{current})}$ and $M^{(\text{new})}$ matrices;

Calculate log-likelihoods, $\log P(X|Z, \sigma_a^{(\text{current})}, \sigma_x)$, with $\sigma_a^{(\text{current})}$;

Calculate log-likelihoods, $\log P(X|Z, \sigma_a^{(\text{new})}, \sigma_x)$, with $\sigma_a^{(\text{new})}$;

Calculate log-priors, $\log \left(e^{-\frac{1}{2} \frac{\sigma_a^{(\text{current})^2}}{\text{prior variance}}} \right)$, for $\sigma_a^{(\text{current})}$;

Calculate log-priors, $\log \left(e^{-\frac{1}{2} \frac{\sigma_a^{(\text{new})^2}}{\text{prior variance}}} \right)$, for $\sigma_a^{(\text{new})}$;

Compute log acceptance ratio $\log \left(\frac{\text{new_log_likelihood} \times \text{new_log_prior}}{\text{current_log_likelihood} \times \text{current_log_prior}} \right)$;

Sample from uniform distribution;

if $\log(\text{uniform}()) < \log(\text{acceptance ratio})$ **then**

 | return $\sigma_a^{(\text{new})}$;

else

 | return $\sigma_a^{(\text{current})}$;

end

Updating σ_x and σ_a within the Gibbs Sampler

During each iteration of the Gibbs Sampler, the hyperparameters σ_x and σ_a are updated. This is achieved by calling the Metropolis-Hastings update functions described in the previous slides.

- The proposal variance for σ_x is calculated as $0.1 \times \sigma_x$.
- The proposal variance for σ_a is calculated as $0.1 \times \sigma_a$.
- The function `metropolis_step_sigma_x()` updates σ_x based on the calculated proposal variance and the prior variance.
- Similarly, `metropolis_step_sigma_a()` updates σ_a .

These update steps are a crucial part of ensuring the Gibbs Sampler explores the parameter space effectively.

The Beta-Bernoulli model: sampling matrix A

We add the update of the matrix A:

Algorithm 8: GibbsSamplerbetabernoulli

Data: $X, A, \sigma_x, \sigma_a, \text{BurnInIterations}, \text{Iterations}$

Result: collection of Z matrices

```
for iter  $\leftarrow$  1 to BurnInIterations + Iterations do
  for i  $\leftarrow$  1 to N do
    K  $\leftarrow$  number of non-empty columns of  $Z_{-i}$ ;
    for j  $\leftarrow$  1 to K do
      | Sample already observed features;
    end
    | Sample new features;
  end
  Update  $\sigma_x$ ;
  Update  $\sigma_a$ ;
  Update A ;
  Compute  $E[A|X, Z]$  and  $\log[P(X, Z)] \leftarrow \text{Eq. (12)+(21)}$ 
  if iter > BurnInIterations then
    | Add Z to the collection;
```

Gaussian Prior for A

We then generalized this model w.r.t. different possible priors of A.
At each iteration of the Gibbs Sampler, we update A and sample from it.

Let $\tau = \frac{1}{\sigma_{A_{n+1}}^2}$:

Algorithm 9: Update A

Data: Z, X, σ_x, σ_A

Result: updated A

$$\tau = \frac{1}{\sigma_{A_n}^2} I_k + \frac{1}{\sigma_x^2} Z^T Z;$$

$$\sigma_{A_{n+1}}^2 = \frac{1}{\tau}; \quad // \text{posterior variance of A};$$

$$\mu_{A_{n+1}} = \sigma_{A_{n+1}}^2 Z^T X \frac{1}{\sigma_x^2}; \quad // \text{posterior mean of A};$$

for $k \leftarrow 0, 1, \dots, K$ **do**

for $d \leftarrow 0, 1, \dots, D$ **do**

$A_{k,d} \sim \mathcal{N}(\mu_{A_{n+1}}(k, d), \sigma_{A_{n+1}}(k, k))$ // updated values of A

end

end

Alternative Prior for A

Now we consider the following prior for the i -th row of A :

$$A_i \mid \mu_i, \sigma^2 \sim \mathcal{N}(\mu_i, \sigma^2),$$

with $\mu_i \sim \mathcal{N}(0, c \sigma^2)$ and $\sigma^2 \sim \text{invgamma}(a, b)$

Algorithm 10: Update A and σ_A

Data: Z, X, a, b, c

Result: updated A

Update a and b ;

$\sigma^2 \sim \text{invgamma}(a, b)$; // posterior variance

for $d \leftarrow 0, 1, \dots, D$ **do**

$\mu_d \sim \mathcal{N}(0, c \sigma^2)$; // posterior mean, with c : multiplicative constant

end

for $k \leftarrow 0, 1, \dots, K$ **do**

for $d \leftarrow 0, 1, \dots, D$ **do**

$A_{k,d} \sim \mathcal{N}(\mu_d, \sigma^2)$; // updated values of A

end

end

- Aggiungere qua le domande
- Come scegliere i valori iniziali di a e b
- Come scegliere c
- In generale, è corretto usare Metropolis-Hastings per l'aggiornamento di σ_x e σ_a , oppure è possibile fare sampling dalle full-conditionals $P(\sigma_x|\text{resto})$ e $P(\sigma_a|\text{resto})$ e fare update con Gibbs sampler? Oppure posso ricavare una distribuzione nota della posterior dal prodotto $\text{likelihood} \times \text{prior}$?
- Se Metropolis-Hastings è appropriato, assumo come prior su σ_x e σ_a una normale con media zero e varianza 1?
- Come metodo per generare la proposal sigma, è corretto quello indicato nell'algoritmo?